

Assignment #3

DUE: December 2nd before class (1:00 PM)

Late assignments will be penalized at the rate of 10% per day

Note there is a **hard cutoff for submission** of this assignment **on December 8, 2022**

Total points: 100 possible

Please complete two problems for this assignment: problem 1 (Difference-in-Differences) and any one of the remaining two problems. I recommend choosing to explore a method that aligns with your research interests, particularly for future planned projects.

Updated November 11: You may ignore parts (e) through (h) of Problem 1—please complete parts (a) through (d) only.

You may use either R or STATA to complete this assignment. Please submit a compiled version of your code that includes the output. In R, this can be achieved by “Compile Report” button in RStudio; in STATA this can be done by creating a .log file. On the first line of your code after the .log or has been started (if using STATA) please also include the following:

In R: `name <- Sys.info() name[7]`
In STATA: `display "`c(username)'"`

Recall that when it comes to grading, commenting your code is your friend. It helps you organize your thoughts and communicates to me more about what you are trying to do. Unless indicated, a complete answer to each part of the assignment will include either a regression table or a figure—both should well-formatted and easily readable—and text describing and interpreting what you are presenting.

1. **Difference-in-Differences: COVID-19 and Sourdough Consumption.** (Adapted from Nick Huntington-Klein and Peter Nencka.) During the early days of COVID-19, there was a brief craze for homemade sourdough bread, as stores were out of yeast (sourdough can be made at home using yeast from the air and does not require store-bought yeast). We will be estimating whether COVID lockdowns actually increased interest in sourdough bread.

We will be measuring interest in sourdough bread using Google Trends data in the USA. Google Trends tracks the popularity of different search terms over time. We will be comparing the popularity of the search term "sourdough" against the control groups: the search terms "cereal," "soup," and "sandwich," the popularity of which we suspect might not have been meaningfully affected by COVID lockdowns.

- a. Load the data set "a3_p1_sourdough_trends.csv" and look through the data. Make a line graph with `date` on the x-axis and `hits` on the y-axis, with a separate line for each `keyword`. Also add a vertical line for the "start of the pandemic" which we'll decide for our purposes is March 15, 2020. Describe what you find – does it lend support for a particular hypothesis? That is, what is (a) the shape any potential effect takes (i.e. is it a permanent increase in popularity? Temporary?), and (b) whether you might be concerned about any of the control groups we've chosen.
- b. Create a "Treated" indicator for sourdough. Perform a test of whether the prior trends differ between the treated and control groups, using a linear trend. Then, if you were concerned about any of the control groups in part (b), drop them (and keep them dropped for the rest of the assignment) and rerun the test. Interpret your test results—does anything concern you here?
- c. Create a `relative_month` variable by shifting the `date` variable back 15 days (so that the treatment day is the first day of the month) and then taking the month of the resulting date. Also create an `After` variable equal to 1/0 (or True/False) if the date is March 15 or afterwards.
 - i. Illustrate how the values of `relative_month` you get line up with `date`, and subtract a number from `relative_month` so that the last period just before treatment (Feb 16-Mar 14) is 0. (Also, change the Jan 1-14 month so it's one less than the Jan 15-Feb 14 month).
- d. Then, use two-way fixed effects to estimate the difference-in-difference estimate of the effect of lockdown on sourdough popularity. What are your two fixed effects? Report and interpret your results, with your standard errors clustered at the `keyword` level.
- e. Now, let's allow the effect to be dynamic over time. Estimate a difference-in-difference model allowing the effect to differ by relative month (using `relative_month = 0` as a reference period), with standard errors clustered at the keyword level. Show the results graphically and interpret them. Does anything about your findings concern you?
- f. Perform an analysis of the weights used in your dynamic TWFE regression. How does this affect your interpretation in (e)?
- g. Briefly describe an estimand that you would be interested in for this research question. Then select an appropriately robust TWFE estimator, and implement it. Have your results changed? How?
- h. What do you conclude about the effect of the COVID-19 pandemic on sourdough bread consumption?

2. **Synthetic Control.** Many countries and provinces have enacted programs to incentivize organ donation, but what are the effects of these programs? Bilgel and Galle (*Journal of Health Economics*, 2015) found strong support for the implementation of tax deductions on the volume of living (nonrelated) organ donations. More recently, many have argued that transitioning from a tax *deduction* to a tax *credit* is a more equitable way to handle incentives for living organ donation.

In this problem, we will assess the effects of one such transition, when Louisiana started a tax *credit* for donations in 2015. We will use US data on kidney donations from the Organ Procurement and Transplantation Network (OPTN).

- a. Load the “a3_p2_organ_donations.csv” dataset. Present a summary statistics table for three variables: donations (total), donations (living), and waiting list additions. Also present a figure showing the rates of living organ donations for each state over time, with a vertical line at 2015, the year of Louisiana’s adoption – make sure that Louisiana is clearly visible. Discuss your results. Note that all variables should be **population-adjusted**.
- b. Would a typical difference-in-differences strategy here suffice to evaluate Louisiana’s policy? Why/why not? Why would/wouldn’t a synthetic control approach be better?
- c. Based on (a) and your poking around in the data, are there any states that should be dropped from your donor pool? Justify your decision, and drop any states you select for the rest of the assignment.
- d. Construct a synthetic control using your donor pool and the following variables: all pre-treatment levels of the outcome, pre-treatment levels of total donations, state population, and state GDP. Report the relative weights of both variables and donor units. Discuss what you see.
- e. Show a balance table between synthetic and real Louisiana. Also include the main figure showing the trends for real and synthetic Louisiana. Was your synthetic control construction successful? Justify, and speculate as to what might have contributed to this.
- f. What is the estimated effect? Show graphically and discuss.
- g. Conduct two placebo tests:
 - i. First, rerun the synthetic control with each state in your donor pool as “treated” at 2015. Show the estimated effect for Louisiana in the full distribution of these estimated placebo effects. What is the percentile of Louisiana’s estimated MSPE compared to those in the full distribution?
 - ii. Second, rerun the synthetic control with the actual date of implementation in 2018 (the date the law went into effect). How does this change your estimated effect?
- h. What do you conclude about Louisiana’s decision to move from a tax deduction to a tax credit? What are the policy implications of your findings?

- 3. Quantile and Nonparametric Regression.** In the U.S., traditional Medicare is a universal, publicly-funded health insurance that is available only to individuals aged 65 and older. This problem assesses how qualifying for Medicare (by turning 65) gives individuals new from medical expenditure risk, and how that differs across the population. This problem is based on [Barcellos and Jacobson, 2015](#).

At age 65, out-of-pocket expenditures drop by 33 percent at the mean and 53 percent at the ninety-fifth percentile. Medical-related financial strain, such as difficulty paying bills and collections agency contact, is dramatically reduced. Nonetheless, using a stylized expected utility framework, the gain from reducing out-of-pocket expenditures accounts for only 18 percent of the social costs of financing Medicare. This calculation ignores any direct health benefits from Medicare or any indirect health effects due to reductions in financial stress.

- a. Load the “a3_p3_medicare.RData” dataset. What is the distribution of out-of-pocket medical spending (“totexp”) for individuals before they are on Medicare versus after? Does this tell you anything about the financial protection afforded to Medicare patients – why or why not?
- b. What is a naïve OLS estimate of the effect of having Medicare (i.e., age ≥ 65) on expected total expenditures? Control for any variables you feel are relevant, and be sure to use the “perwt” weighting variable (throughout this problem).
- c. Now report quantile regression results at 20 evenly spaced quantiles of the income distribution (“faminc”). Report the raw regression coefficients in a figure (feel free to drop outliers, as long as that is noted). How do you interpret these results?
- d. Scale these coefficients by the quantile value of the income distribution and recreate the figure. How does this scaling change the interpretation of the results?
- e. Instead of quantile regression coefficients, construct a binscatter that illustrates the relationship between the income distribution and overall health expenses. Report two relationships: one for individuals on Medicare and one for individuals who are too young to qualify. Adjust for the relevant covariates and weights you used in parts (b) and (c). How does the difference between these two curves change across the distribution? What is the interpretation of that change?
- f. What do your results from (e) imply about the use of nonparametric regression here? Can you specify a framework to recover a nonparametric conditional density estimator of the effect of Medicare coverage over the income distribution? Discuss the assumptions you would need and the procedure you would implement for this approach to be valid. Note that you do not have to implement it, necessarily, but you are welcome to if you feel up to it.
- g. What do you conclude about the distributional effects of Medicare as a form of risk protection? Can you think of any practical mechanisms (open back-doors) that explain your findings? What are the policy implications of your findings?